

DRISHTI

See beyond your resume.

An AI-powered career intelligence agent exploring whether a candidate's future direction can be better understood through experience, interests, projects, learning activities, and career goals — combined.

Ushasree Jakilinki · AI Product & Agentic AI Portfolio

Drishti (Sanskrit): vision, insight, perspective

The Thesis

“Career decisions should not be based solely on who a person has been — but also on who they are becoming.”

Traditional platforms ask

“What jobs match my resume?” — matching against historical keywords, so a career in transition stays invisible to the machine.

Drishti asks

“What jobs match who I am becoming?” — and ultimately: “what do I need to do to get there?”

Drishti = Resume (Past) + Activity (Present) + Aspirations (Future) → Opportunities + Gaps

The Problem

The people who need discovery most — candidates in transition — are served worst.

Keyword matching sees only the past

Recommendation engines match historical resume data. Career progression, transferable skills, active learning, and aspirations are invisible — so candidates manually review dozens of irrelevant postings.

A lived example

A candidate's resume says Product Manager. Meanwhile she is building AI agents, publishing product case studies, studying agentic AI and Responsible AI, and targeting AI Product Management. Every platform keeps recommending... Product Manager roles. The evidence of her transition lives outside the resume — in projects, courses, GitHub, and stated goals — where no engine looks.

Result: high effort, low efficiency, and transitions penalized instead of recognized.

The Insight — Candidate DNA

Three signal types. Not all career signals have equal value — and none should be ignored.

PAST

Historical Signals

Evidence of proven capability

Resume · roles · education · certifications · industries

v1 weight: 0.4

NOW

Current Signals

Evidence of active skill development

Projects · courses · GitHub · portfolio · learning activities

v1 weight: 0.3

NEXT

Intent Signals

Evidence of future direction

Target roles · career goals · interests · desired technologies

v1 weight: 0.3

Weighted together into a single Drishti Score — the weights themselves are a tunable experiment, refined by user feedback.

Drishti v1 — A Deliberately Thin Slice

Prove the thesis end-to-end before adding anything. One user, local data, no APIs.

IN — v1

- Resume upload & parsing (PDF/DOCX) with a user validation step
- Signals form: courses, projects, links, target roles, goals
- Career Trajectory Profile (JSON data contract)
- Curated local dataset: 100–300 JDs in 1–2 target domains
- Past / Present / Future scoring → Drishti Score
- Explanations across all three signals + gap analysis
- Feedback: “Would you apply?” Yes / No / Maybe


OUT — parked, not abandoned

- Live job APIs (USAJobs, Adzuna, Jooble, RemoteOK, Remote)
- Multi-agent orchestration (Phase 3, when it becomes the lesson)
- Automated applications & resume rewriting
- Career roadmap generation & progress tracking
- Multi-user / SaaS platform
- Enterprise integrations

Backlog discipline: ideas are documented, not built.

Scope-creep is a named risk — observed personally during discovery. The backlog absorbs it.

v1 Architecture — Modules First, Agents Later



The data contract

The Trajectory Profile JSON schema is frozen in Phase A — every module (and every future agent) speaks it. Modules get promoted to true agents in Phase 3, when orchestration itself becomes the learning objective.

Local-first by principle

Ollama (Qwen/Llama/Mistral) + sentence-transformer embeddings + SQLite + ChromaDB + Streamlit. Resumes and goals never leave the machine. Zero cloud costs.

What a Drishti Recommendation Looks Like

AI Product Manager — recommended

Drishti Score: 0.86

✓ Fits your PAST

Product management and delivery leadership; enterprise stakeholder experience; CSPO.

✓ Fits your PRESENT

Building AI agents and case studies; studying agentic AI, Responsible AI, and evaluation frameworks.

✓ Fits your FUTURE

Stated target: AI Product Management and agentic AI solution roles.

⚠ Gaps to close

Shipped AI product at scale not yet demonstrated · certification preferred by employer · consider one deployed end-to-end demo. (Gaps feed the future roadmap engine: what to learn, build, and show next.)

Every recommendation explains itself across all three signals — no unexplained scores, ever.

Evaluation — Same Discipline, New Domain

“It works” and “it’s trustworthy” are different milestones.

Extraction accuracy

Parsed profile vs. hand-labeled resume fields; user validation step catches the rest

Recommendation quality

Human eval per job: Would you apply? Target: $\geq 70\%$ Yes/Maybe

Hallucination testing

No fake jobs, no fake skills, no fake certifications — outputs validated against sources

Performance

End-to-end recommendations in under 10 seconds on local data

Feedback (Yes/No/Maybe) is both the quality metric and the future training signal for re-weighting the Drishti Score.

Risks & Mitigations

Resume parsing inaccuracies

Mandatory user review/edit of the extracted profile before matching.

Recommendation hallucinations

Explainability + confidence scoring; dedicated hallucination test suite.


Limited job data (no APIs in v1)

Curated 100–300 JD dataset, scoped to 1–2 domains done well.

Scope creep during discovery

Observed personally. Ideas go to the backlog; build validates few hypotheses at a time.

Roadmap — Thin Slices, Each One Shippable



Future phases — the Career Intelligence vision

Live job APIs → multi-agent orchestration (Resume, Preference, Discovery, Matching, Insight, Feedback agents) → gap analysis, learning recommendations, and personalized career roadmaps: “what do I need to do to get there?”

DRISHTI

See beyond your resume.

Traditional systems ask: “What jobs match my resume?”

Drishti asks: “What jobs match who I am becoming — and what do I need to do to get there?”

Ushasree Jakilinki

linkedin.com/in/ushasreejakilinki · github.com/auj012